

AI Data Access Regulation Policy

MT&T Corporation - Effective Date: January 1, 2026

Classification: INTERNAL - CONFIDENTIAL

Document Version: 2.1 | Last Updated: March 15, 2026

1. Purpose and Scope

This policy establishes the framework for regulating access to data used by artificial intelligence systems within MT&T Corporation. It applies to all employees, contractors, partners, and automated systems that access, process, or generate data through AI/ML platforms deployed on the Databricks Lakehouse Platform.

The policy covers all workspaces, model serving endpoints, AI Gateway configurations, foundation model APIs, vector search indexes, and Genie rooms operating within MT&T's Databricks environment.

2. Data Classification for AI Systems

All data accessed by AI systems must be classified into one of four tiers: (a) Public - Data that can be freely used for AI training and inference, (b) Internal - Data restricted to internal AI applications with standard access controls, (c) Confidential - Data requiring enhanced controls including audit logging, encryption at rest and in transit, and role-based access, (d) Restricted - Data that requires explicit approval from the Data Governance Council before AI processing, including PII, PHI, financial records, and customer communications.

AI models trained on Confidential or Restricted data must be tagged in Unity Catalog with appropriate sensitivity labels. Model serving endpoints exposing such models must implement AI Gateway guardrails to prevent data leakage in responses.

3. Access Control Requirements

All AI platform access must use Unity Catalog-managed identities. Service principals must be used for automated workflows and model serving endpoints. Personal access tokens are prohibited for production AI workloads.

Access to model serving endpoints must follow the principle of least privilege. Each endpoint must have an explicit access control list (ACL) defining which users, groups, and service principals can invoke it. Wildcard permissions (ALL USERS) are prohibited for endpoints serving Confidential or Restricted data.

Cross-workspace access to AI resources requires approval from both the source and destination workspace administrators. All cross-workspace AI data flows must be documented in the AI Data Flow Registry maintained in the `cmegdemos_catalog.ai_governance` schema.

4. Audit and Monitoring Requirements

All AI data access events must be captured in the `system.access.audit` table. The AI Governance team must review audit logs weekly for anomalous access patterns, unauthorized access attempts, and policy violations.

Model serving endpoint usage must be monitored via `system.serving.endpoint_usage`. Alerts must be configured for: (a) More than 100 failed requests per hour from a single requester, (b) Access from unrecognized IP addresses, (c) Token consumption exceeding 10x the 30-day average, (d) Any access to endpoints during maintenance windows.

The AI Governance Dashboard application must be used as the primary monitoring interface. All stakeholders must have access to the dashboard and Genie room for self-service query capabilities.

5. Data Retention and Deletion

AI inference logs containing input/output data must be retained for 90 days for audit purposes and then automatically purged. Model training datasets must follow the retention schedule defined by MT&T's Enterprise Data Retention Policy (EDRP-2025).

When a model serving endpoint is decommissioned, all associated inference tables, vector search indexes, and cached embeddings must be deleted within 30 days. The AI Governance team must verify deletion through the `v_underutilized_endpoints` view.

6. Compliance and Enforcement

Violations of this policy will be reported to the Chief Data Officer and may result in immediate revocation of AI platform access. Repeat violations may result in disciplinary action as outlined in MT&T's Code of Conduct.

Annual compliance audits will be conducted by the Internal Audit team using the AI Access Audit view (`v_ai_access_audit`) to verify adherence to access control requirements. Findings must be remediated within 30 business days.