

AI Model Governance and Lifecycle Policy

MT&T Corporation - Effective Date: January 1, 2026

Classification: INTERNAL - CONFIDENTIAL

Document Version: 2.1 | Last Updated: March 15, 2026

1. Purpose

This policy defines the governance framework for the complete lifecycle of AI models within MT&T Corporation, from development through deployment, monitoring, and retirement. It ensures all AI models meet quality, security, fairness, and compliance standards before and during production deployment.

2. Model Development Standards

All AI models must be developed using MT&T-approved frameworks and registered in Unity Catalog's Model Registry. Each model must have a Model Card documenting: purpose and intended use cases, training data sources and lineage, performance metrics and evaluation results, known limitations and biases, responsible AI assessment results.

Foundation models accessed via the Databricks Foundation Model API or AI Gateway must be pre-approved by the AI Ethics Committee. Currently approved foundation models include: Meta Llama 3.3 70B, Anthropic Claude family, OpenAI GPT-4 family, Google Gemini family, and DBRX. Provisioned throughput endpoints must be justified with a cost-benefit analysis showing expected usage above 60% utilization.

3. Model Deployment Approval Process

Before any model is deployed to a serving endpoint, it must pass through a three-stage approval gate: Stage 1 (Technical Review) - The MLOps team validates model performance, resource requirements, and serving configuration. Stage 2 (Security Review) - The Security team verifies access controls, data classification compliance, and AI Gateway guardrail configuration. Stage 3 (Business Review) - The business owner confirms the model serves an approved use case and budget allocation exists.

Emergency deployments may bypass Stage 3 with approval from two VP-level sponsors, but must complete full review within 5 business days post-deployment.

4. Production Monitoring Requirements

All production model serving endpoints must be monitored using the AI Governance Dashboard. The following metrics must be tracked: Request volume and trend via `v_serving_endpoint_daily`, Token consumption and cost via `v_ai_cost_daily`, Error rates and latency via `v_ai_gateway_daily`, User access patterns via `v_ai_access_audit`.

Model owners must configure alerts for: Error rate exceeding 5% over a 1-hour window, Latency exceeding SLA thresholds (p99 < 2 seconds for chat, < 500ms for embeddings), Cost exceeding 120% of monthly budget allocation, Drift in input distribution detected via Databricks

Lakehouse Monitoring.

5. Endpoint Utilization and Cost Management

The AI Governance team must review endpoint utilization monthly using the `v_underutilized_endpoints` view. Endpoints classified as 'Idle' (zero requests in 30 days) must be decommissioned within 14 days unless a business justification is provided. Endpoints classified as 'Very Low' (fewer than 100 requests in 30 days) must provide a utilization improvement plan within 30 days.

Provisioned throughput endpoints with utilization below 40% must be converted to pay-per-token or scaled down. The estimated cost savings must be reported to the FinOps team monthly.

6. Model Retirement Process

Models must be retired when they are replaced by a newer version, no longer serve an active business use case, or fail to meet performance or compliance standards. The retirement process includes: 30-day deprecation notice to all consumers, traffic migration to replacement endpoint, archival of model artifacts and inference logs, deletion of serving endpoint and associated resources, confirmation via the AI Governance Dashboard.

7. Responsible AI Requirements

All customer-facing AI models must undergo bias testing before deployment. Models used for decisions affecting customers (credit scoring, service eligibility, pricing) must provide explainability features. AI Gateway guardrails must be configured to block harmful content, PII leakage, and off-topic responses for all chat-based endpoints.