

AI Cost Management and FinOps Policy

MT&T Corporation - Effective Date: January 1, 2026

Classification: INTERNAL - CONFIDENTIAL

Document Version: 2.1 | Last Updated: March 15, 2026

1. Purpose

This policy establishes the financial governance framework for AI/ML workloads within MT&T Corporation. It ensures cost transparency, accountability, and optimization across all AI platform resources including model serving, foundation model APIs, model training, and AI-powered analytics.

2. AI Cost Categories and Ownership

AI costs are categorized into five pillars, each with designated cost owners: (a) Foundation Model Serving - Costs for Anthropic, OpenAI, and Gemini model serving endpoints, owned by the AI Platform team. (b) Custom Model Inference - Costs for real-time inference endpoints serving custom ML models, owned by respective business unit ML teams. (c) Model Training - GPU compute costs for fine-tuning and training, owned by the Data Science Center of Excellence. (d) AI-Powered Analytics - Costs for Genie rooms, AI/BI dashboards, and assistant usage, owned by the Analytics team. (e) AI Infrastructure - Vector search, feature serving, and monitoring costs, owned by the MLOps team.

Each cost category is tracked via the `v_ai_cost_daily` view using billing SKU mappings: `ENTERPRISE_ANTHROPIC_MODEL_SERVING`, `ENTERPRISE_OPENAI_MODEL_SERVING`, `ENTERPRISE_GEMINI_MODEL_SERVING`, `ENTERPRISE_MODEL_TRAINING_*`, and `ENTERPRISE_SERVERLESS_REAL_TIME_INFERENCE_*`.

3. Budget Allocation and Approval Thresholds

Annual AI budgets are allocated at the business unit level and tracked monthly. Budget thresholds for AI resource provisioning: (a) Under \$5,000/month - Team lead approval, (b) \$5,000-\$25,000/month - Director approval with cost justification, (c) \$25,000-\$100,000/month - VP approval with ROI analysis, (d) Over \$100,000/month - C-suite approval with strategic business case.

Provisioned throughput endpoints require pre-approval based on expected monthly cost. Pay-per-token endpoints are preferred for variable or unpredictable workloads.

4. Cost Monitoring and Alerting

The AI Governance Dashboard's Cost Observatory tab must be reviewed weekly by cost owners. Automated alerts must be configured for: Daily cost exceeding 150% of the 30-day rolling average, Monthly cost on track to exceed 110% of budget allocation, Any single endpoint consuming more than 25% of the business unit's AI budget, New AI SKUs appearing in billing that were not in the approved budget.

The v_ai_cost_daily view provides the granular data needed for cost analysis. Cost owners must create Genie room questions for their specific cost monitoring needs.

5. Cost Optimization Requirements

Quarterly cost optimization reviews must be conducted by each business unit. Required optimization actions: Consolidate underutilized endpoints (utilization_tier = 'Idle' or 'Very Low'), Right-size provisioned throughput based on actual usage patterns, Implement request caching for repeated queries to foundation models, Use smaller models (e.g., Llama 8B instead of 70B) where quality requirements allow, Batch inference for non-real-time workloads to reduce per-request overhead.

The AI FinOps team must report monthly savings achieved through optimization to the CFO. Target: 15% year-over-year cost efficiency improvement while maintaining service quality.

6. Chargeback and Showback

All AI costs must be attributed to business units using Unity Catalog tags and usage_context fields in model serving requests. Teams must populate the usage_context field with department, project, and environment identifiers for all endpoint invocations.

Monthly chargeback reports are generated from the v_ai_cost_daily and v_serving_endpoint_daily views and distributed to budget owners by the 5th business day of each month.