

AI Platform Acceptable Use Policy

MT&T Corporation - Effective Date: January 1, 2026

Classification: INTERNAL - CONFIDENTIAL

Document Version: 2.1 | Last Updated: March 15, 2026

1. Purpose

This policy defines acceptable and prohibited uses of MT&T Corporation's AI platform resources. It applies to all users of the Databricks-based AI infrastructure including model serving endpoints, foundation model APIs, AI Gateway, Genie rooms, and associated tools.

2. Acceptable Uses

MT&T's AI platform is approved for the following use cases: Customer experience improvement through intelligent chatbots and virtual assistants, Network optimization and predictive maintenance using ML models, Fraud detection and security analytics, Revenue analytics and customer churn prediction, Internal knowledge management and document analysis, Data quality monitoring and automated anomaly detection, Self-service analytics through Genie rooms.

3. Prohibited Uses

The following uses of MT&T's AI platform are strictly prohibited: Generating or processing illegal content, Using AI for employment decisions without HR and Legal approval, Accessing competitor data or performing unauthorized competitive intelligence, Training models on customer data without appropriate consent and privacy review, Using foundation models to generate code for unauthorized systems, Sharing model serving endpoint credentials with external parties, Bypassing AI Gateway guardrails or rate limits, Using production AI resources for personal projects or non-MT&T purposes.

4. User Responsibilities

All AI platform users must: Complete MT&T's AI Ethics and Governance training before accessing AI resources (renewed annually), Report suspected policy violations to the AI Governance team within 24 hours, Use the usage_context field in model serving requests to identify their department and project, Review and acknowledge this policy annually through the MT&T compliance portal.

Model serving endpoint owners must: Monitor their endpoints through the AI Governance Dashboard, Respond to utilization alerts within 2 business days, Participate in quarterly cost optimization reviews, Maintain up-to-date Model Cards for all deployed models.

5. Monitoring and Enforcement

Compliance with this policy is monitored through: The AI Governance Dashboard and its underlying system table views, Automated alerts on the v_ai_access_audit view for policy violations, Quarterly usage reviews by the AI Governance Committee, Annual audits by Internal

Audit using all governance views.

Policy violations are classified as: Minor - First-time infractions with no data exposure (verbal warning), Moderate - Repeated infractions or potential data exposure (written warning, mandatory retraining), Severe - Confirmed data breach, unauthorized access, or prohibited use (access revocation, disciplinary action), Critical - Willful policy violation causing material harm (termination, legal action).

6. Policy Review and Updates

This policy is reviewed quarterly by the AI Governance Committee, which includes representatives from Legal, Security, Data Science, MLOps, and Business Operations. Updates are communicated through MT&T's policy management portal and require 30 days notice before taking effect.

Questions about this policy should be directed to the AI Governance team or asked via the AI Governance Q&A Genie room.